

```
from transformers import pipeline, set_seed, GPT2LMHeadModel, GPT2Tokenizer
import torch

# Reproducibility for exact generations
set_seed(42)

generator = pipeline('text-generation', model='gpt2',
                     device=0 if torch.cuda.is_available() else -1)

prompt = "Once upon a time in a distant galaxy,"
outputs = generator(prompt, max_new_tokens=100, num_return_sequences=3,
                    temperature=0.8, top_p=0.95, do_sample=True, pad_token_id=50256)

print("=== Text Generations from Pipeline ===")
for i, output in enumerate(outputs):
    print(f"\nGeneration {i+1}:\n{output['generated_text']}")
    print("-" * 80)

# Manual approach with more control
tokenizer = GPT2Tokenizer.from_pretrained('gpt2')
model = GPT2LMHeadModel.from_pretrained('gpt2')

# Move the model to the correct device to avoid RuntimeError
model.to(generator.device)

tokenizer.pad_token = tokenizer.eos_token

prompt2 = "In a world where AI has become conscious,"
inputs = tokenizer(prompt2, return_tensors="pt").to(generator.device)

outputs2 = model.generate(**inputs, max_new_tokens=120,
                          temperature=0.7, top_p=0.9, do_sample=True, repetition_penalty=1.2)

print("\n=== Manual Generation with Full Control ===")
print(tokenizer.decode(outputs2[0], skip_special_tokens=True))
```

```
/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:93: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
warnings.warn(
config.json: 100% 665/665 [00:00<00:00, 16.7kB/s]
Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and fast
WARNING:huggingface_hub.utils._http:Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN t
model.safetensors: 100% 548M/548M [00:04<00:00, 184MB/s]

Loading weights: 100% 148/148 [00:00<00:00, 327.28it/s, Materializing param=transformer.wte.weight]

GPT2LMHeadModel LOAD REPORT from: gpt2
Key | Status | |
-----+-----+--
h.{0...11}.attn.bias | UNEXPECTED | |

Notes:
- UNEXPECTED :can be ignored when loading from different task/architecture; not ok if you expect identical arch.
generation_config.json: 100% 124/124 [00:00<00:00, 2.89kB/s]

tokenizer_config.json: 100% 26.0/26.0 [00:00<00:00, 990B/s]

vocab.json: 1.04M/? [00:00<00:00, 10.5MB/s]
merges.txt: 456k/? [00:00<00:00, 9.45MB/s]
tokenizer.json: 1.36M/? [00:00<00:00, 18.1MB/s]

Passing `generation_config` together with generation-related arguments=({'pad_token_id', 'num_return_sequences', 'do_sample'
Both `max_new_tokens` (=100) and `max_length` (=50) seem to have been set. `max_new_tokens` will take precedence. Please refe
=== Text Generations from Pipeline ===

Generation 1:
Once upon a time in a distant galaxy, a vast galaxy that had been destroyed, and had been completely destroyed, the world of

Now, as it slowly turned its attention to the next step, the human race had discovered something that had surprised everyone

A small, dark, black hole.

A world in which the humans and the aliens had been separated.

A world in which the humans
-----

Generation 2:
Once upon a time in a distant galaxy, the galactic community had already developed a technological civilization capable of
```